

HBase

Welcome to the world of NoSQL

- Introduction to HBase
- Why to use HBase
- What HBase is not
- Characteristics
- Data Model
- Architecture
- HBase vs HDFS
- HBase vs RDBMS



- HBase is an open source, non-relational, distributed database modeled after Google's BigTable and written in Java.
- Think of it as a sparse, consistent, distributed, multidimensional, sorted map:
 - labeled tables of rows
 - row consist of key-value cells:
 - (row key, column family, column, timestamp) -> value

- Storing large amounts of data (TB/PB)
- Storing unstructured or variable column data
- Big Data with random read and writes
- Horizontally scalable
- Strongly consistent reads and writes
- Simple Java API
- Integration with Map/Reduce framework

What HBase is not

- Not an SQL database
- Not relational
- No joins
- No fancy query language and no sophisticated query engine
- No transactions out-of-the box
- No secondary indices out-of-the box
- Not a drop-in replacement for your RDBMS

- Non-relational (NoSQL)
- Column-Oriented
- Multi-Dimensional
- High Availability
- High Performance
- Open Source

- HBase is a "NoSQL" database.
- "NoSQL" is a general term meaning that the database doesn't support SQL as its primary access language. There are many types of NoSQL databases:
 - **Key-Value Store** – It has a Big Hash Table of keys & values (Example- Riak, Amazon S3 (Dynamo))
 - **Document-based Store** - It stores documents made up of tagged elements. (Example- CouchDB)
 - **Column-based Store** - Each storage block contains data from only one column, (Example- HBase, Cassandra)
 - **Graph-based** - A network database that uses edges and nodes to represent and store data. (Example- Neo4J)

- Every row has a row key (analogous to a primary key)
- Rows are stored by row key for fast lookups
- A table may have 1 or more column families
- Common to have a small number of column families
- A column family can have number of columns
- Each row has a timestamp
- Multiple versions of a row can exist

- Table is a distributed sorted map
 - row key + column family + column + timestamp -> value

Row key	Column key	Timestamp	Cell
Row1	info:aaa	1273516197868	valueA
Row1	info:bbb	1273871824184	valueB
Row1	info:bbb	1273871823022	oldValueB
Row1	info:ccc	1273746289103	valueC
Row2	info:hello	1273878447049	i_am_a_value
Row3	info:	1273616297446	another_value

Sorted by Row key and Column

Column family

Column key can be empty

Timestamp is a long value

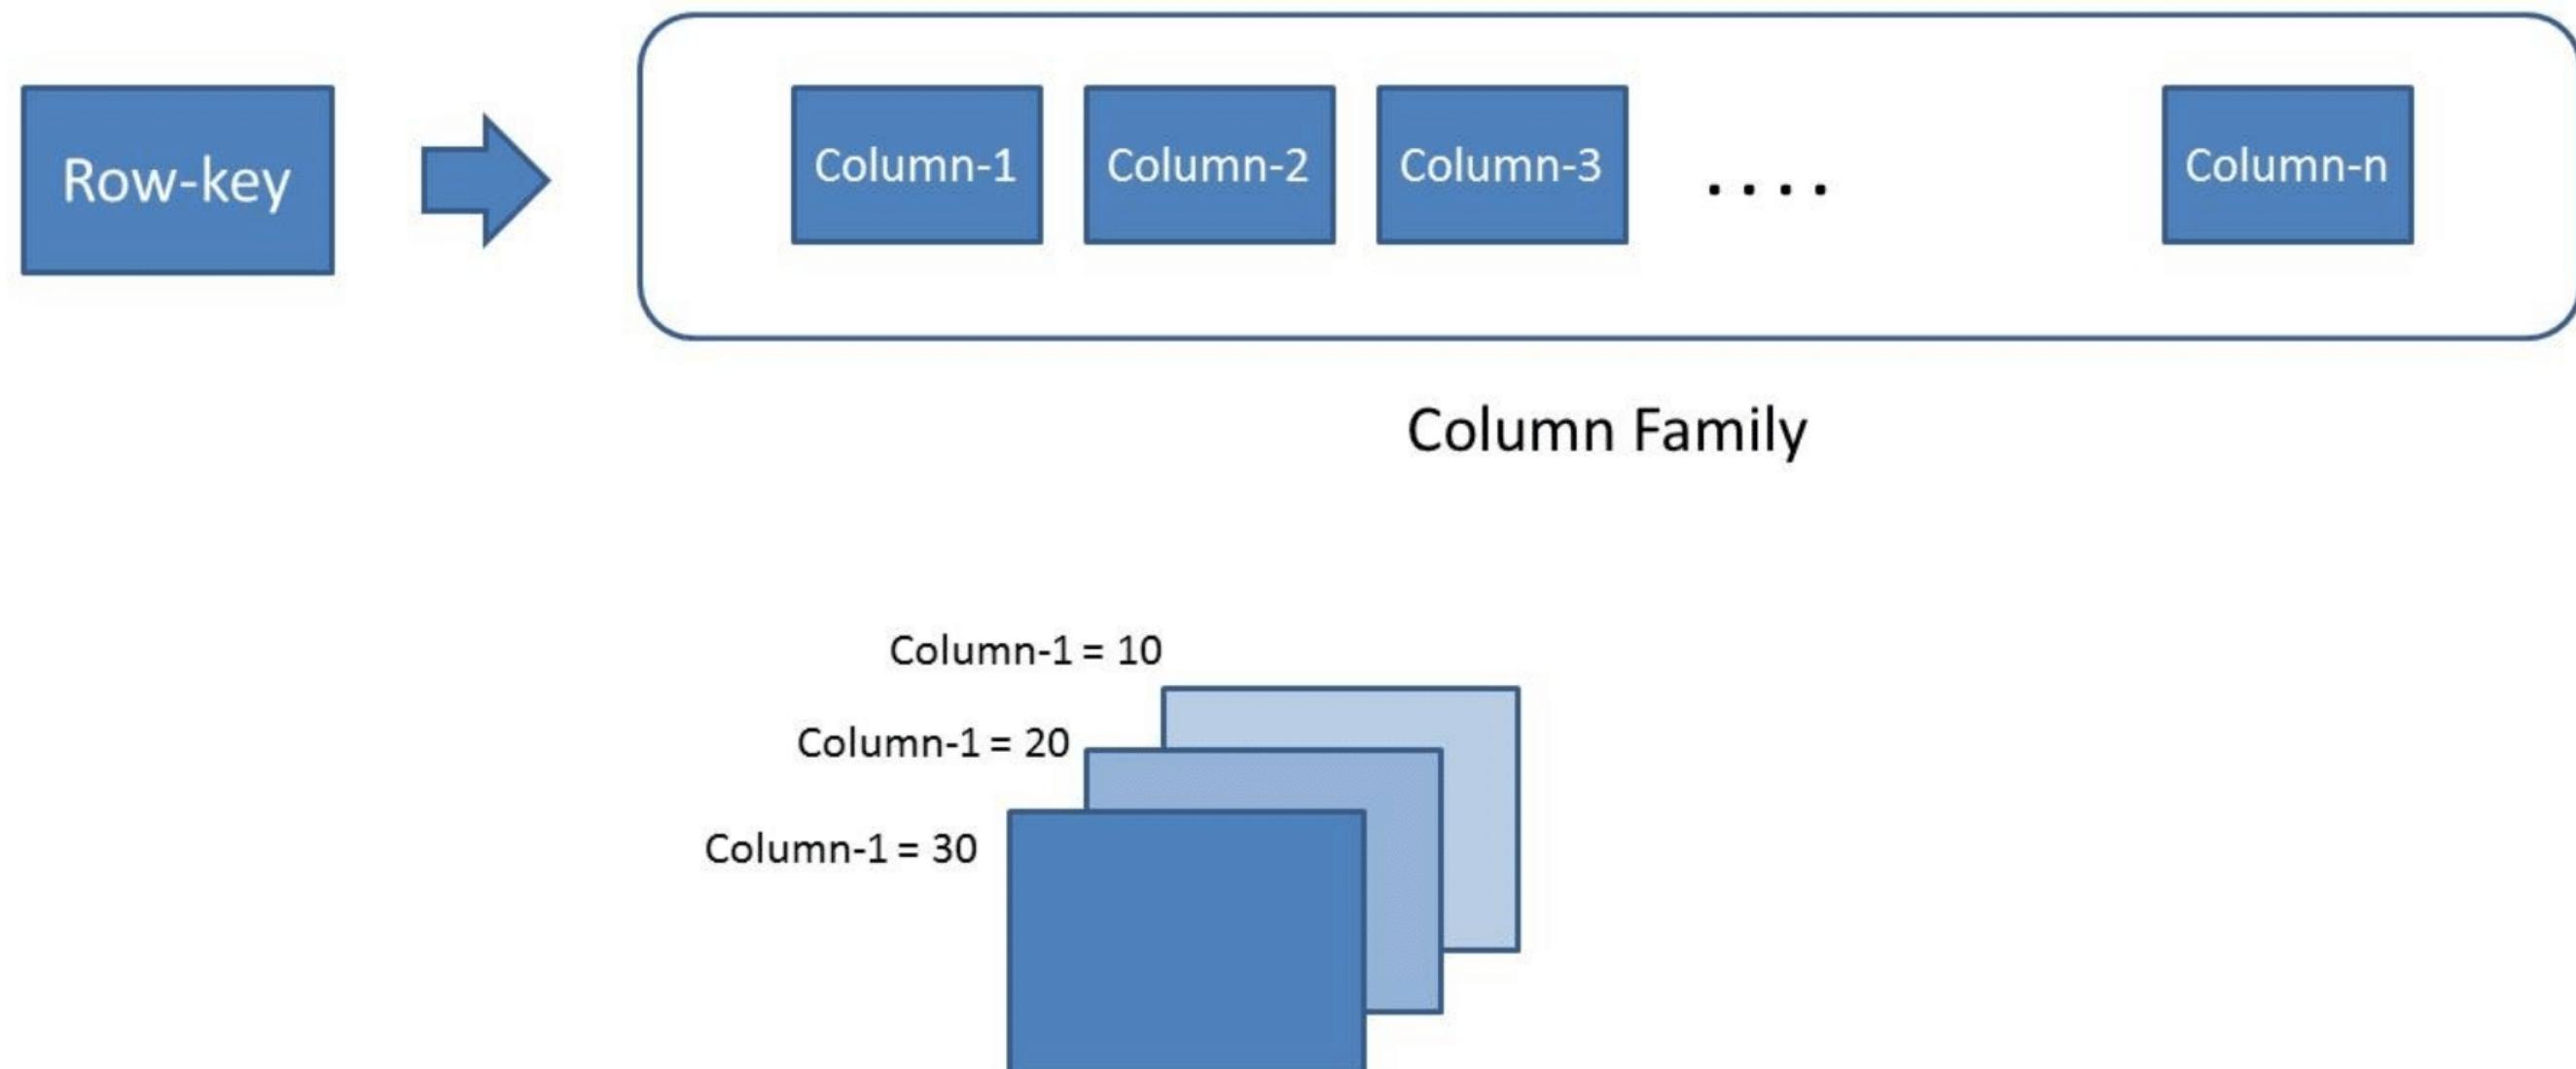
2 Versions of this row

- Stores multiple versions of columns
- With each version a timestamp is stored

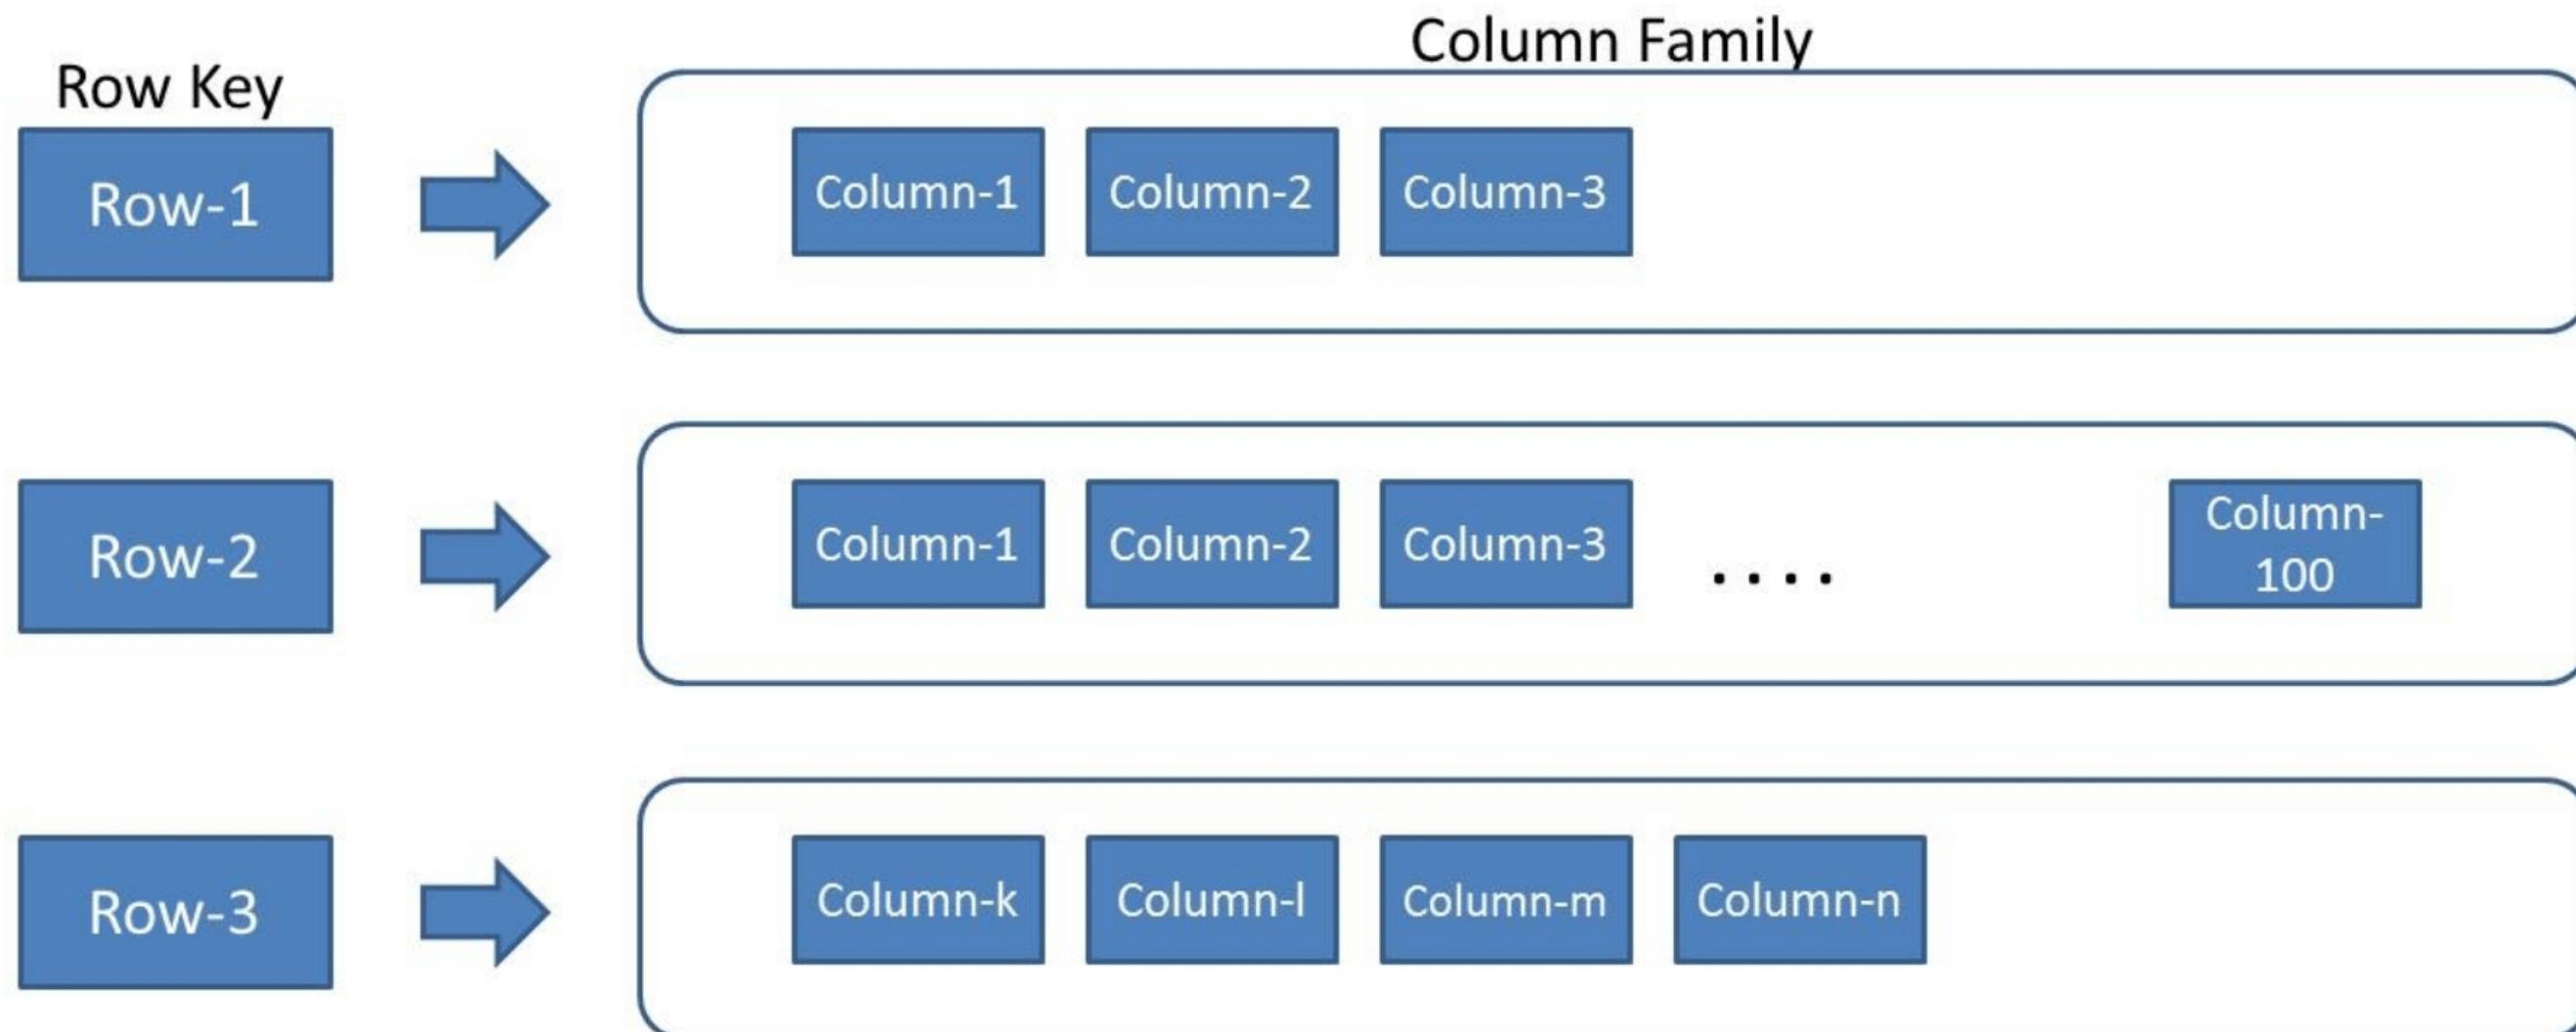
Key	Column	Value	Timestamp
rowA	Fam:foo	New value	1275340679713
rowA	Fam:foo	Old value	1275091706190
rowB	Fam:foo	Some value	1274999316683

Sorted in
descending
order

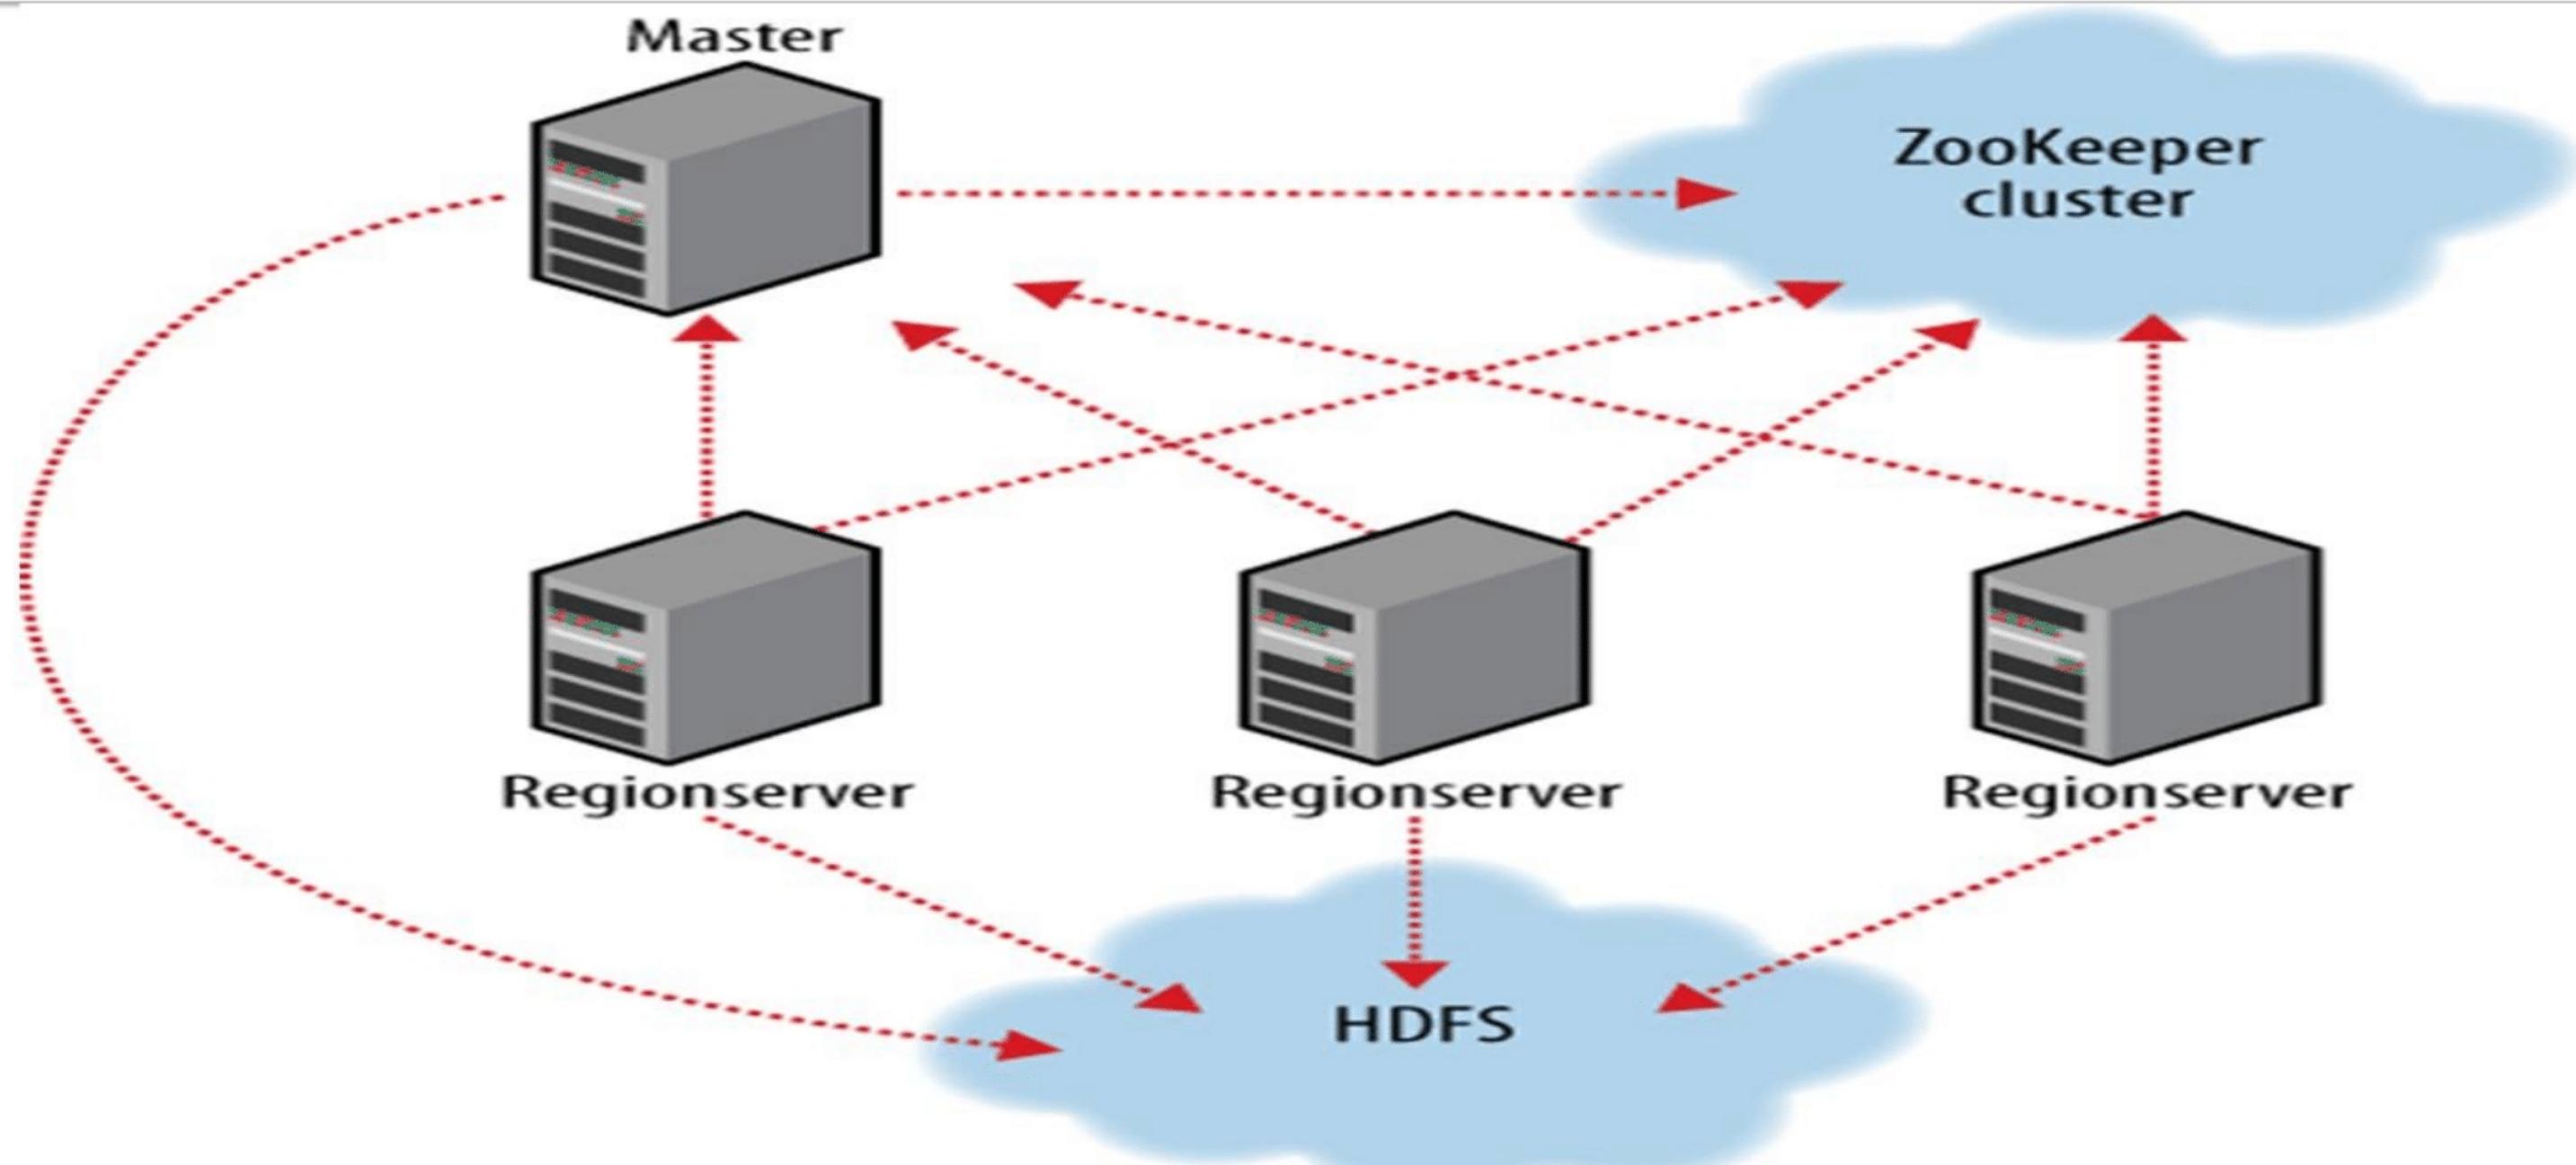
Versions of Data



Create Table specify: Table-Name, Column-Family-Name



HBase Architecture



- Column – oriented database (column families)
- Table consists of Rows, each which has a primary key(row key)
- Each Row may have any number of columns
- Table schema only defines Column familes (column family can have any number of columns)
- Each cell value has a timestamp

- HBase shell
 - Scan, List, Create, Get, Put, Delete
- Java API
 - Get, Put, Delete, Scan,...
- Non-Java Clients
 - Thrift
 - REST
- HBase MapReduce API
 - `hbase.mapreduce.TableMapper;`
 - `hbase.mapreduce.TableReducer;`
- High Level Interface
 - Pig, Hive

- Both are distributed systems that scale to hundreds or thousands of nodes
- HDFS is good for batch processing (scans over big files)
 - Not good for record lookup
 - Not good for incremental addition of small batches
 - Not good for updates
 - Not good for Random read write

- HBase is designed to efficiently address the above points
 - Fast record lookup
 - Support for record-level insertion
 - Support for updates
- HBase updates are done by creating new versions of values

- Facebook Analytics
 - Real-time counters of URLs shared, preferred links
- Twitter
 - 25 TB of message every month
- Mozilla
 - Store crashes report, 2.5 million per day.

