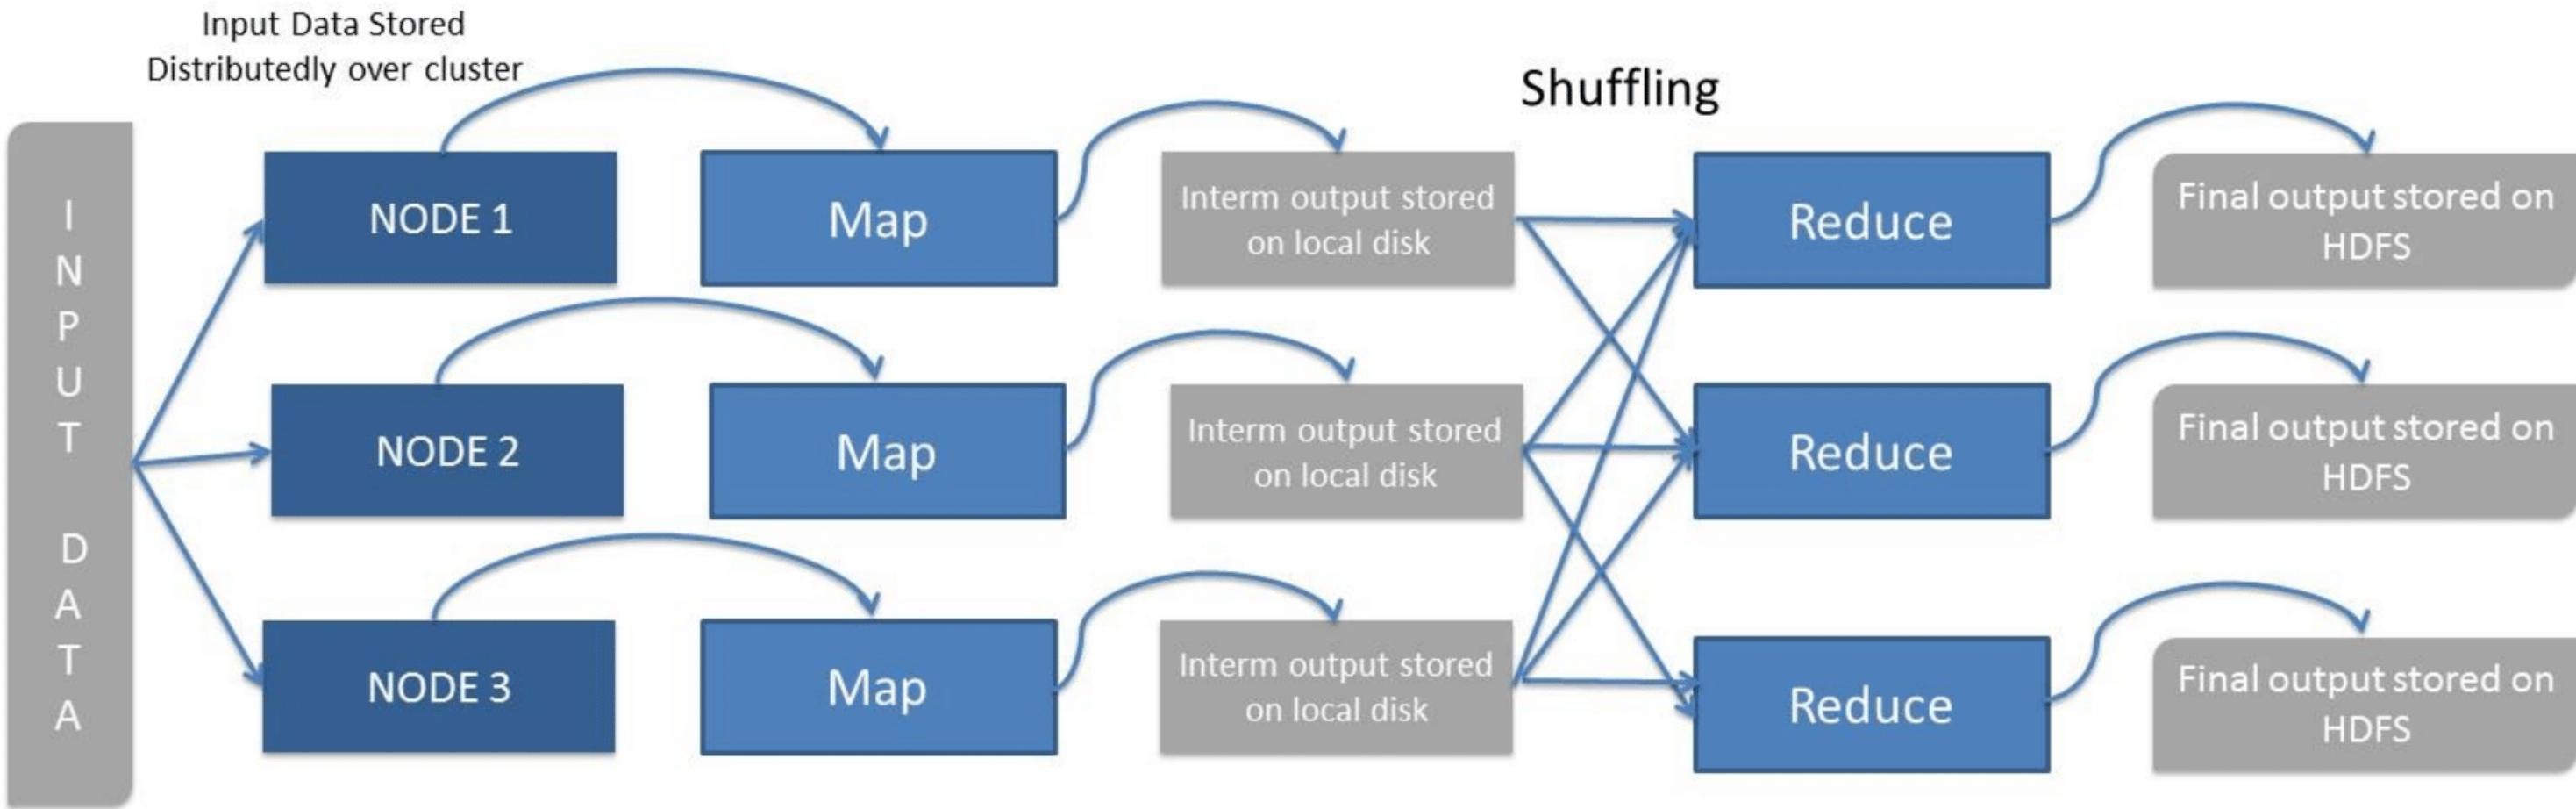


Advanced MapReduce Concepts

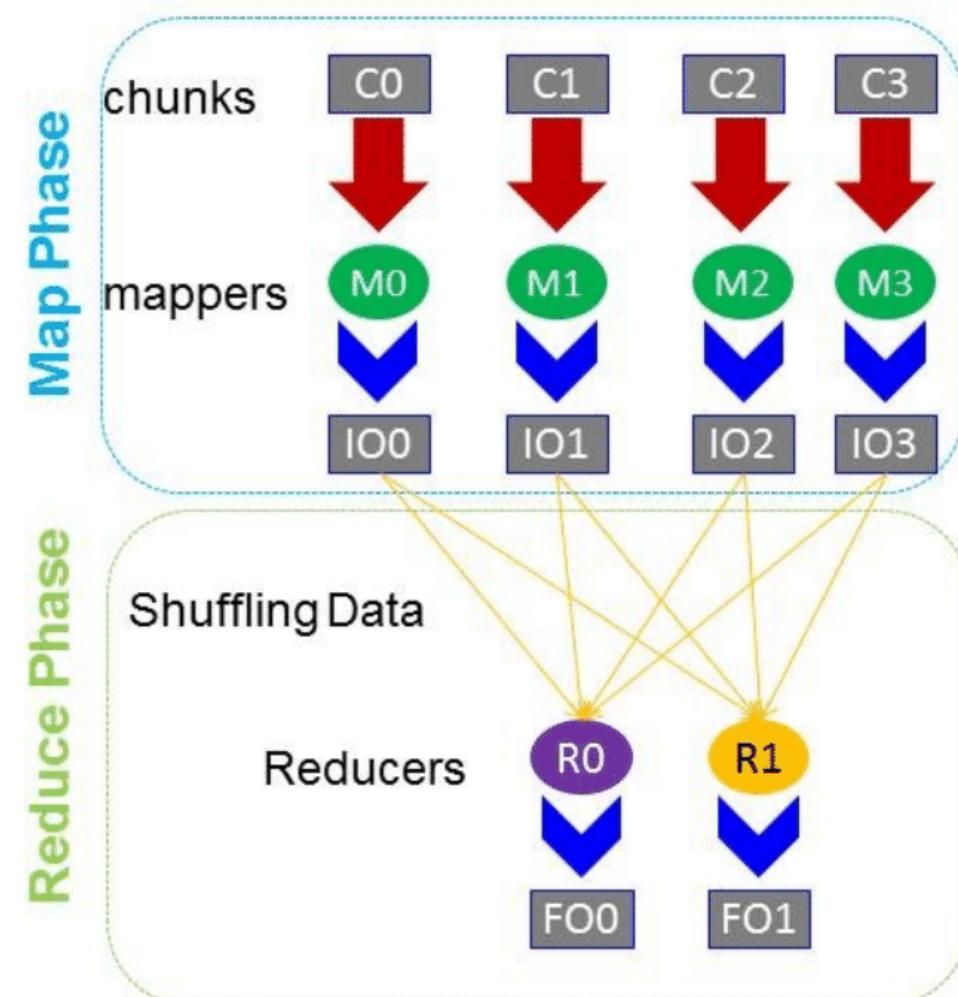
- ✓ Anatomy of MapReduce
- ✓ Hadoop data types
- ✓ Map Reduce – A closer look
- ✓ Map Reduce Components
- ✓ Map Reduce in a Nut-Shell



Anatomy of MapReduce



- In MapReduce, chunks are processed in isolation by tasks called *Mappers*
- The outputs from the mappers are denoted as *intermediate outputs* (IOs)
- The process of bringing together IOs into a set of Reducers is known as *Shuffling and Sort process*
- Intermediate outputs will be given as input into second set of tasks called *Reducers*
- The Reducers produce the *final outputs* (FOs)
- Overall, *MapReduce* breaks the data flow into two phases, map phase and reduce phase



- MR has a defined way of keys and values types for it to move across cluster:
 - Values → Writable
 - Keys → WritableComparable<T>
 - WritableComparable = Writable+Comparable<T>

Java type	Hadoop type
Boolean	BooleanWritable
Byte	ByteWritable
Double	DoubleWritable
Integer	IntWritable
Long	LongWritable
String	Text
null	NullWritable

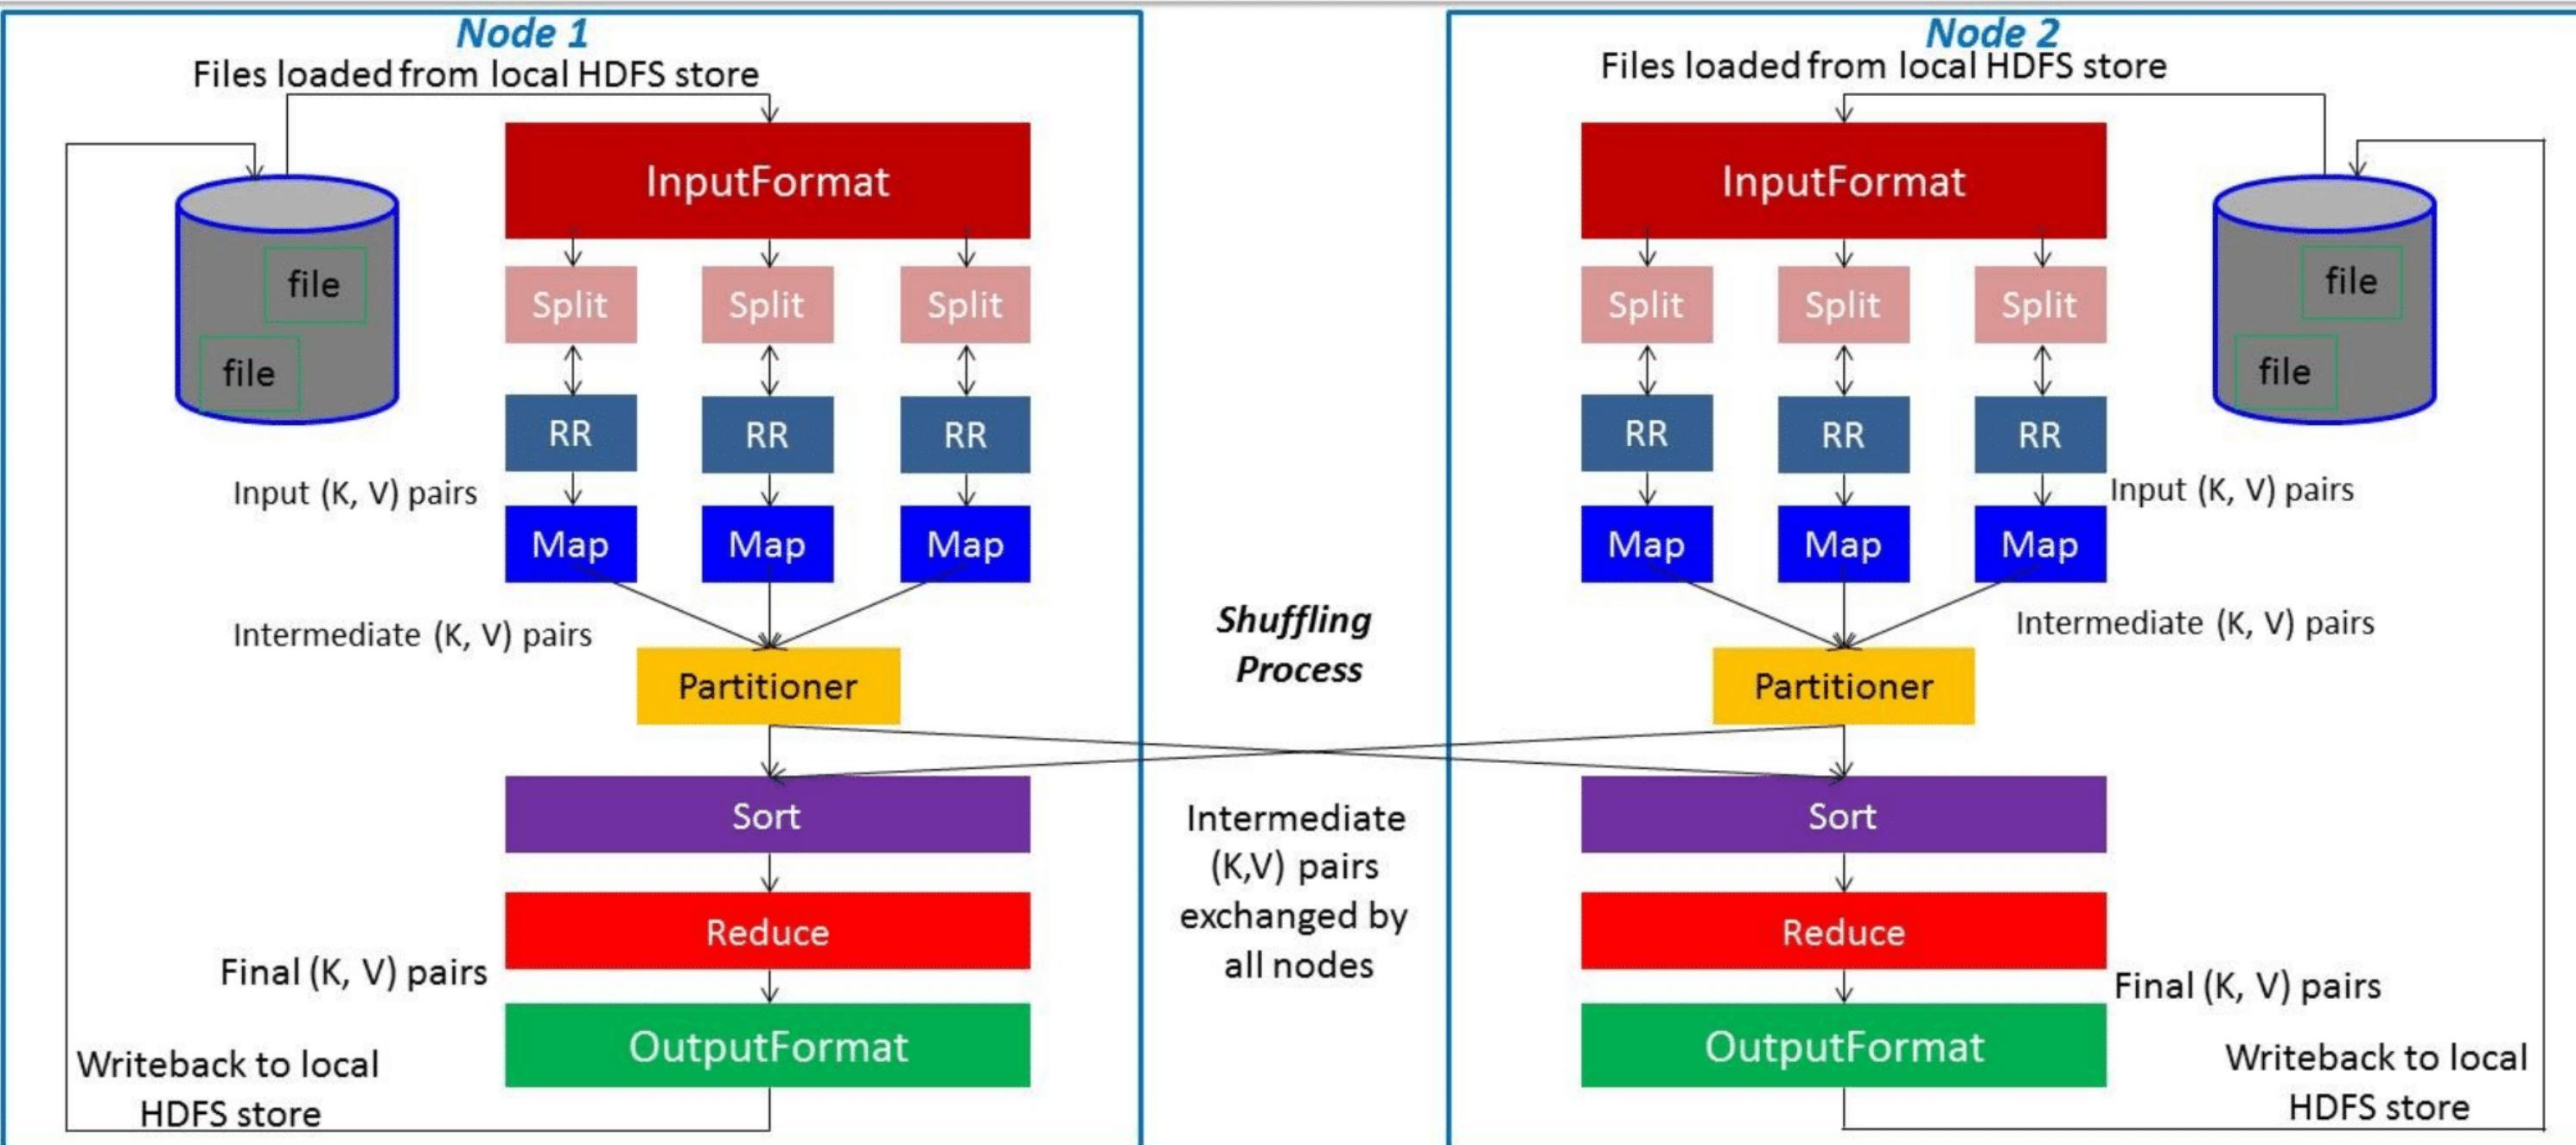
For any class to be value, it has to implement

- org.apache.hadoop.io.Writable
 - write(DataOutput out)
 - readFields(DataInput in)

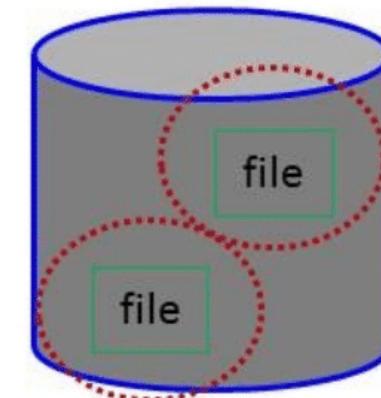
For any class to be key, it has to implement

- org.apache.hadoop.io.WritableComparable<T>
 - write(DataOutput out)
 - readFields(DataInput in)
 - compareTo(T o)

MapReduce – A Closer Look



- Input files are where the data for a MapReduce task is initially stored
- The input files typically reside in a distributed file system (e.g. HDFS)
- The format of input files is arbitrary
 - Line-based log files
 - Binary files
 - Multi-line input records
 - Or something else entirely

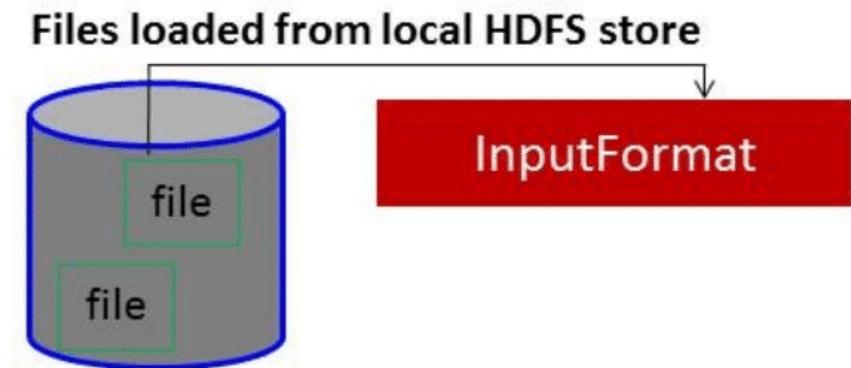


Large File



- Big File is divided into multiple blocks and stored in hdfs.
- This is a physical division of data
- `dfs.block.size` (128MB default size)

- How the input files are split up and read is defined by the *InputFormat*
- *InputFormat* is a class that does the following:
 - Select the files that should be used for input
 - Defines the *InputSplits* that break a file
 - Provides a factory for *RecordReader* objects that read the file



InputFormat (Contd..)

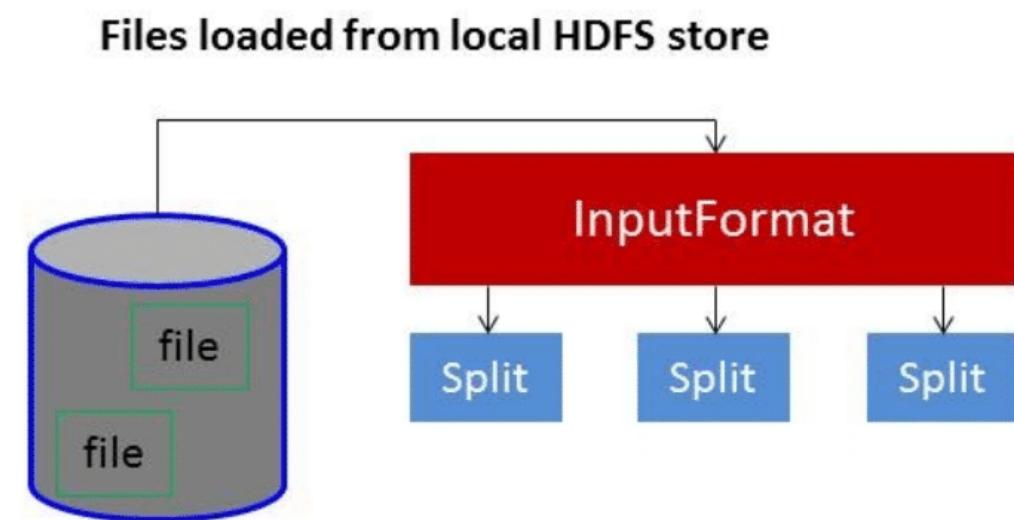
- An InputFormat is responsible for creating the input splits and dividing them into records.
- How we get the data to mapper

```
public abstract class InputFormat<K, V> {  
    public abstract List<InputSplit> getSplits(JobContext context)  
        throws IOException, InterruptedException;  
  
    public abstract RecordReader<K, V>  
        createRecordReader(InputSplit split,  
        TaskAttemptContext context) throws IOException,  
        InterruptedException;  
}
```

The Hadoop framework provides a large variety of input formats:

- **TextInputFormat:** The key is the byte offset, and the value is the line.
- **KeyValueInputFormat:** Key upto first tab char, rest is the value
- **NLineInputFormat:** Similar to TextInputFormat, but it provides n lines of input.
- **MultiFileInputFormat:** An abstract class that lets the user implement an input format that aggregates multiple files into one split.
- **SequenceFileInputFormat:** The input file is a Hadoop sequence file, containing serialized key/value pairs.

- An *input split* describes a unit of work that comprises a single map task in a MapReduce program
- By default, the InputFormat breaks a file up into 64MB splits
- By dividing the file into splits, it allows several map tasks to operate on a single file in parallel
- If the file is very large, this can improve performance significantly through parallelism
- Each map task corresponds to a *single* input split



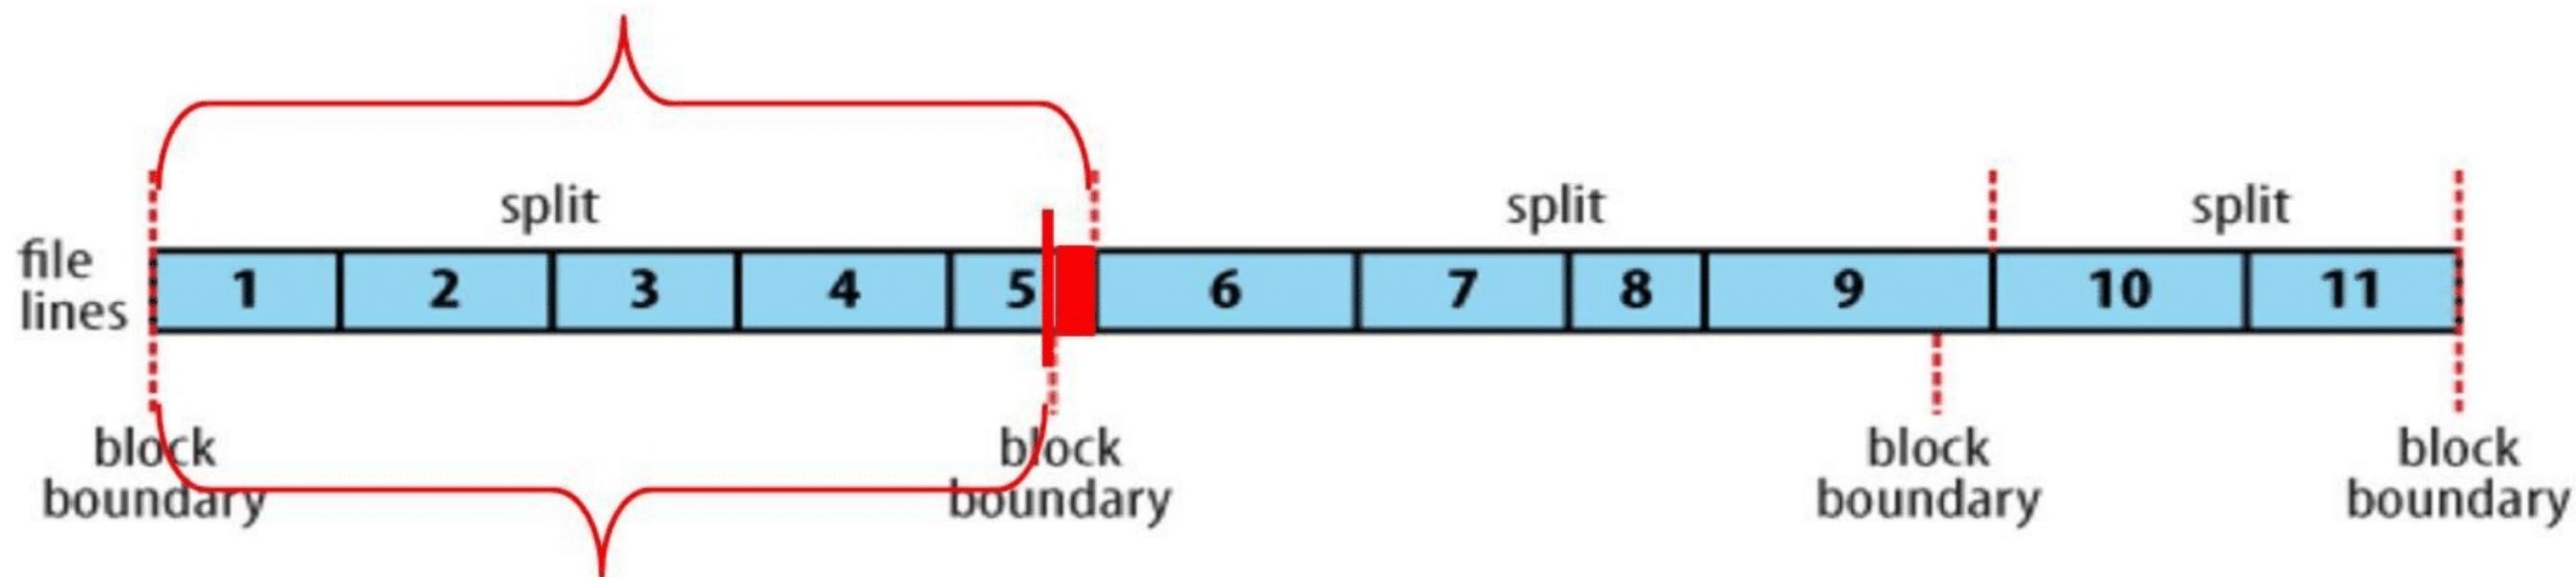
Input split

- A chunk of data processed by a mapper
- Further divided into records
- Map process these records
 - Record = key + value
- How to correlate to a DB table
 - Group of rows → split
 - Row → record

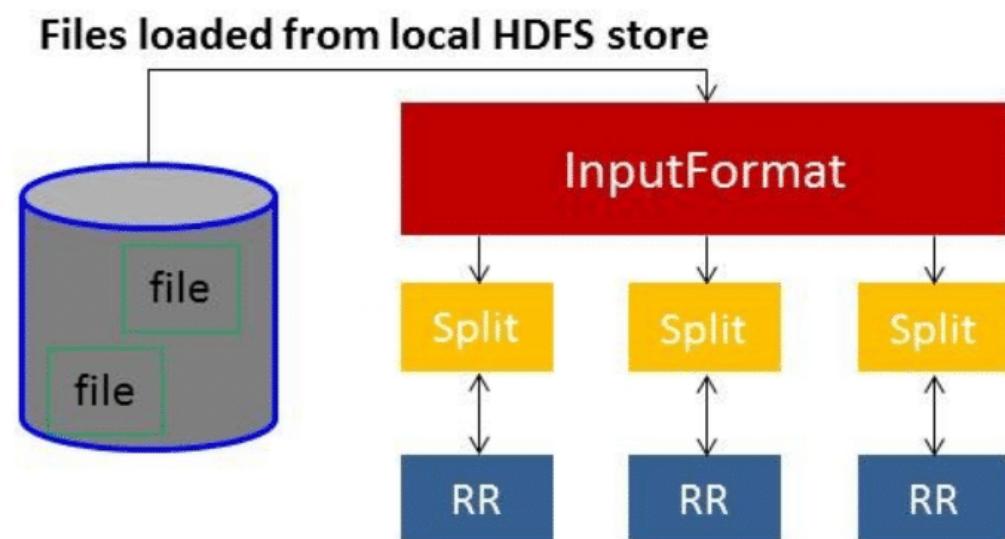


Key	Value
R1	R1
R2	R2
R3	R3

- Split gives logical record boundaries
- Blocks – physical boundaries (size)
- Even in data locality some remote reading is done, a slight overhead.



- The input split defines a slice of work but does not describe how to access it
- The *RecordReader* class actually loads data from its source and converts it into (K, V) pairs suitable for reading by Mappers
- The RecordReader is invoked repeatedly on the input until the entire split is consumed
- Each invocation of the RecordReader leads to another call of the map function defined by the programmer



- Each mapper may emit (K, V) pairs to *any* partition
- Therefore, the map nodes must all agree on where to send different pieces of intermediate data
- The *partitioner* class determines which partition a given (K,V) pair will go to
- The default partitioner computes a *hash value* for a given key and assigns it to a partition based on this result

